

# INTRODUCTION TO DATA MINING

---

Daniel Rodríguez, University of Alcalá

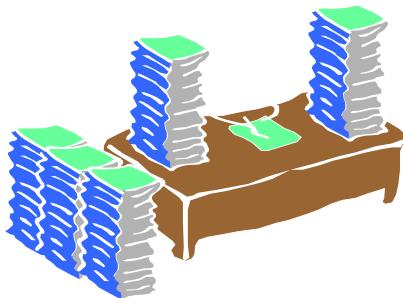


# Outline

- Knowledge Discovery in Datasets
- Model Representation
  - Types of models
    - Supervised
    - Unsupervised
- Evaluation
- (Acknowledgement: Jesús Aguilar-Ruiz for some slides)

# Knowledge Discovery in Datasets

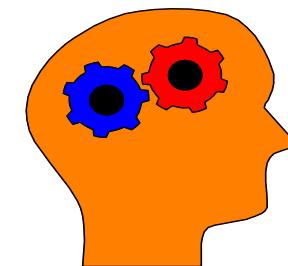
KDD is the automatic extraction of non-obvious, hidden knowledge from large volumes of data.



Large amount  
of data



Data mining  
algorithms?



What knowledge?  
How to represent  
and use it?

# Knowledge Discovery in Datasets

The non-trivial process of identifying valid, novel, potentially useful, and ultimately understandable patterns in data -  
Fayyad, Piatetsky-Shapiro, Smyth (1996)

non-trivial process

valid

novel

useful

understandable

multiple process

Justified patterns/models

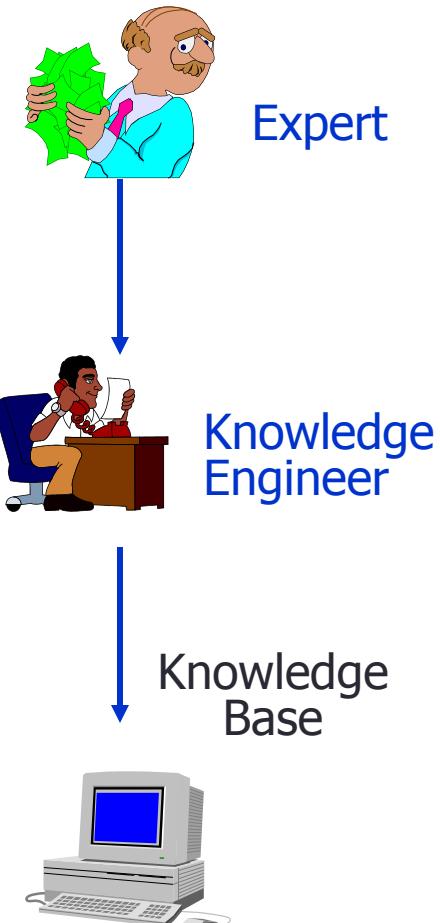
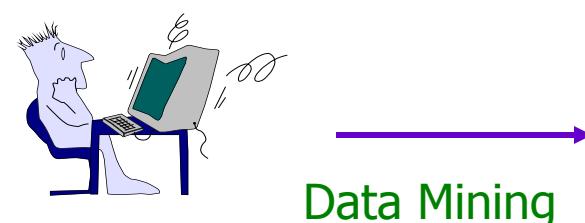
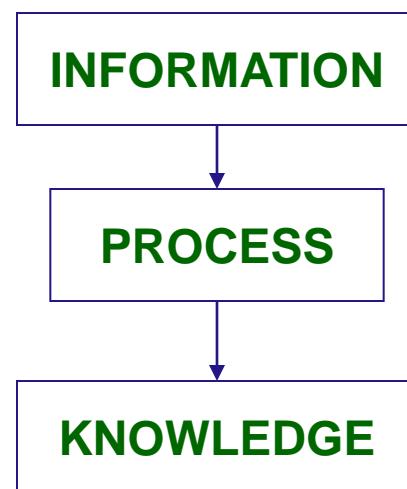
Previously unknown

Can be used

by human and machine

# Knowledge Discovery in Datasets

Impractical manual  
data analysis



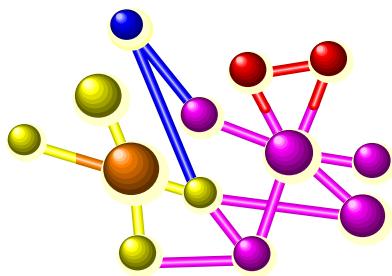
# KDD: Potential Applications

## Business information



- Marketing and sales data analysis
- Investment analysis
- Loan approval
- Fraud detection
- etc.

## Scientific information



- Sky survey cataloging
- Biosequence Databases
- Geosciences: Quakefinder
- etc.

## Manufacturing information

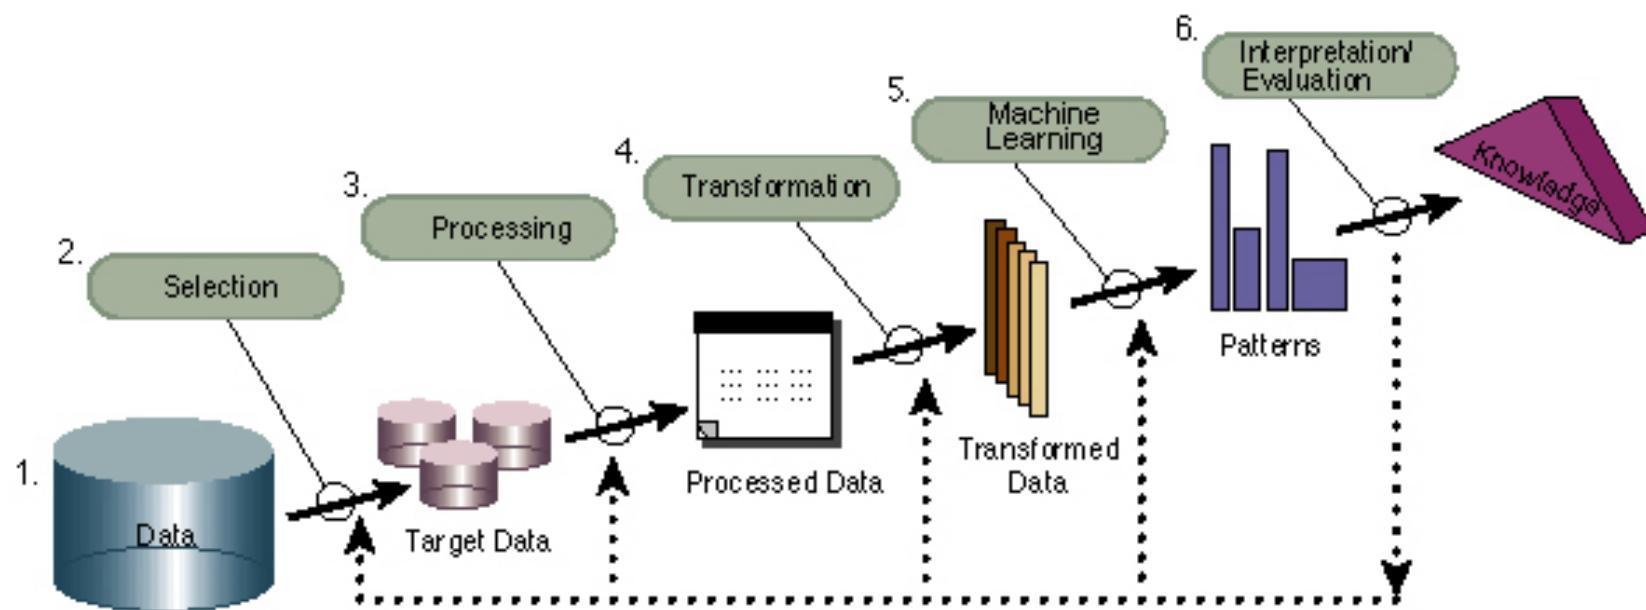


- Controlling and scheduling
- Network management
- Experiment result analysis
- etc.

## Personal information



# KDD Process



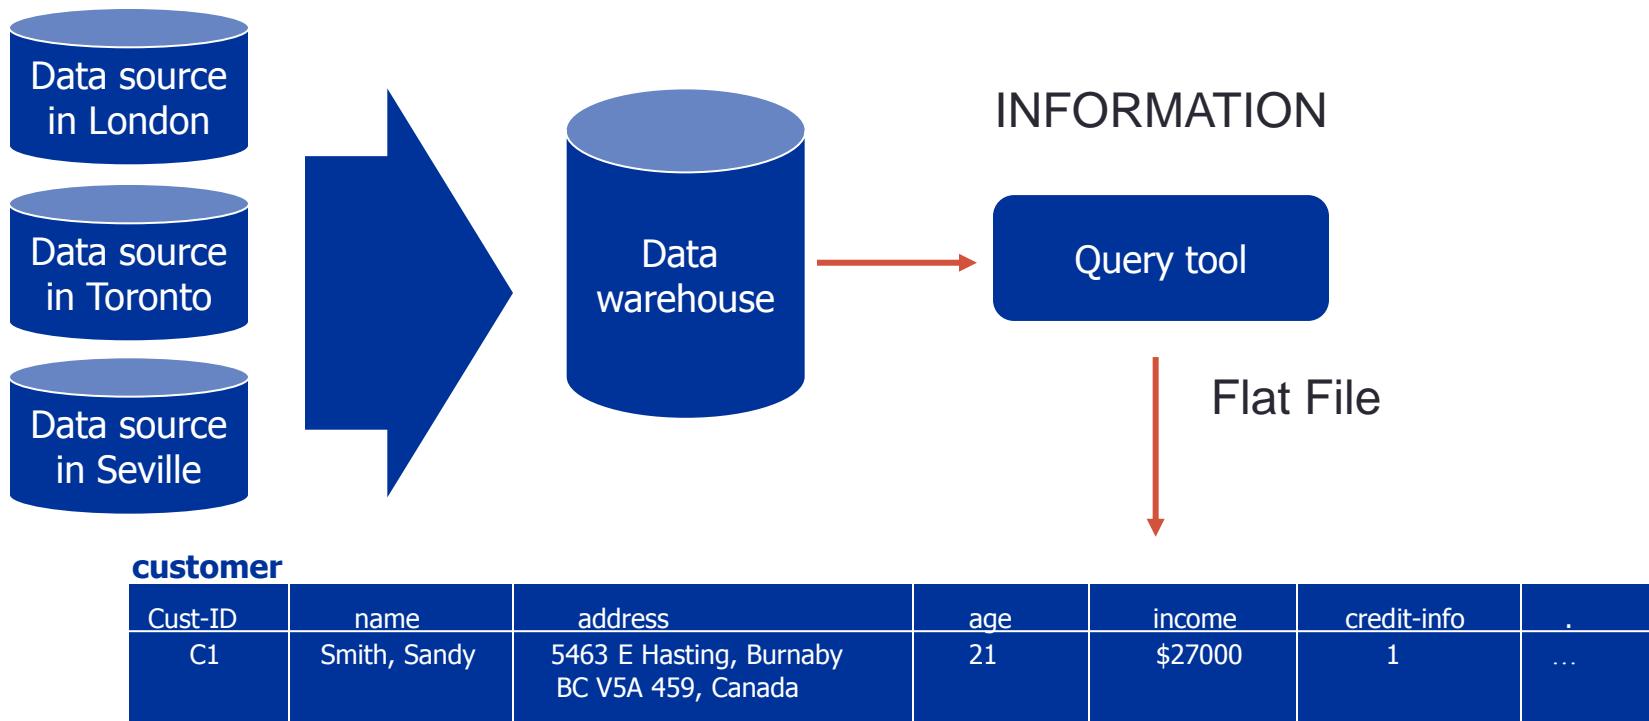
An Overview of the Steps That Compose the KDD Process

# KDD Process

1. Data integration and recopilation
2. Selection, cleaning and transformation
  - Data
3. Machine Learning (data mining)
  - Patterns (e.g., classifiers, rules, etc.)
4. Evaluation and interpretation
  - Knowledge
5. Decision making

# KDD Process: Integration & Recopilation

- Data warehousing, databases, data repositories from different and heterogeneous sources.
- We usually deal with a plain and flat matrix of data



# KDD Process: Selection, cleaning and transformation

- Removing Outliers
- Data Sampling (if there are too much data)
- Missing Values
- Removing redundant and irrelevant attributes (feature selection)
- Derive new attributes from existing ones
  - E.g., Population density from population and area
- Discretisation, normalisation, etc.

# KDD Model Classification

- DM algorithms are traditionally divided into
  - **Supervised learning** which aims to discover knowledge for classification or prediction (predictive)
  - **Unsupervised learning** which refers to the induction to extract interesting knowledge from data (descriptive)
- New approaches also consider **semisupervised learning**: goal is classification but the input contains both unlabeled and labeled data.
  - **Subgroup Discovery** approaches generate descriptive rules are also half way between descriptive and predictive techniques.

# Supervised learning - classifiers

- A classifier resembles a function in the sense that it attaches a value (or a range or a description) to a set of attribute values.
- Induce a classification model, given a database with
  - $m$  instances (samples) characterized by
  - $n$  predicted attributes,  $A_1, \dots, A_n$ ,
  - and the class variable,  $C$  (it is possible to have more than one class).

$A_1$	$\dots$	$A_n$	$C$
$a_{1,1}$	$\dots$	$a_{1,n}$	$c_1$
$\dots$	$\dots$	$\dots$	$\dots$
$a_{m,1}$		$a_{m,n}$	$c_m$

# Supervised Techniques

- **Decision trees** are trees where each leaf indicates a class and internal nodes specifies some test to be carried out.
  - There are many tree-building algorithms such as C4.5 (Quilan) or ID3.
- **Rule induction**
  - If condition then class<sub>label</sub>
  - If ... then ... else if ... (hierarchical)
- **Lazy techniques** store previous instances and search similar ones when performing classification with new instances
  - **k-nearest neighbour** (k-NN) is a method for classifying objects based on closest training examples in the feature space.
- **Numeric prediction**
  - Regression Techniques
- **Neural Networks** are composed by a set of nodes (units, neurons, processing elements) where each node has input and output and performs a simple computation by its node function.
- **Statistical Techniques**
  - **Bayesian networks classifiers** assign a set of attributes  $A_1, A_2, \dots, A_n$  to a class  $C_j$  such that  $P(C_j | A_1, A_2, \dots, A_n)$  is maximal.
- **Meta-techniques**

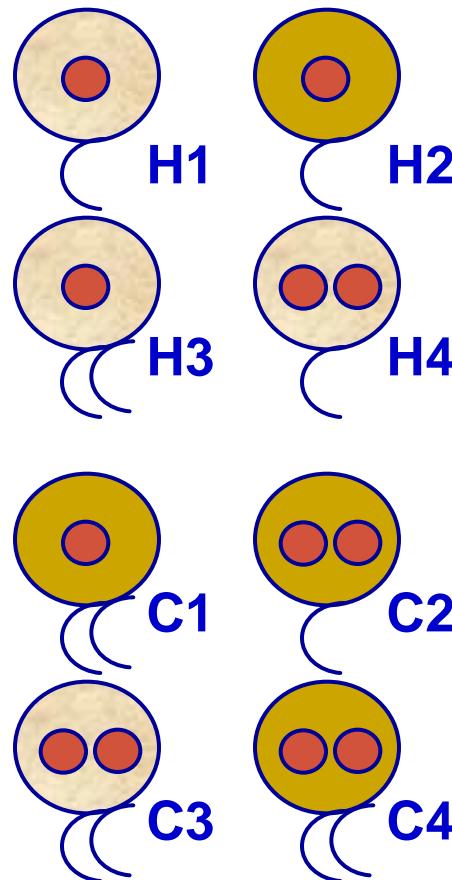
# Unsupervised Techniques

- Association
  - Association Rules, APRIORI
    - E.g., rules among supermarket items
- Clustering
  - Tree clustering: join together objects (e.g., animals) into successively larger clusters, using some measure of similarity or distance.
  - Algorithms: K-Means, EM (Expectation Maximization)
- There is no ‘class’ attribute or class considered as another attribute

$A_1$	...	$A_n$
$a_{1,1}$	...	$a_{1,n}$
...	...	...
$a_{m,1}$		$a_{m,1}$

# KDD Representation (Models)

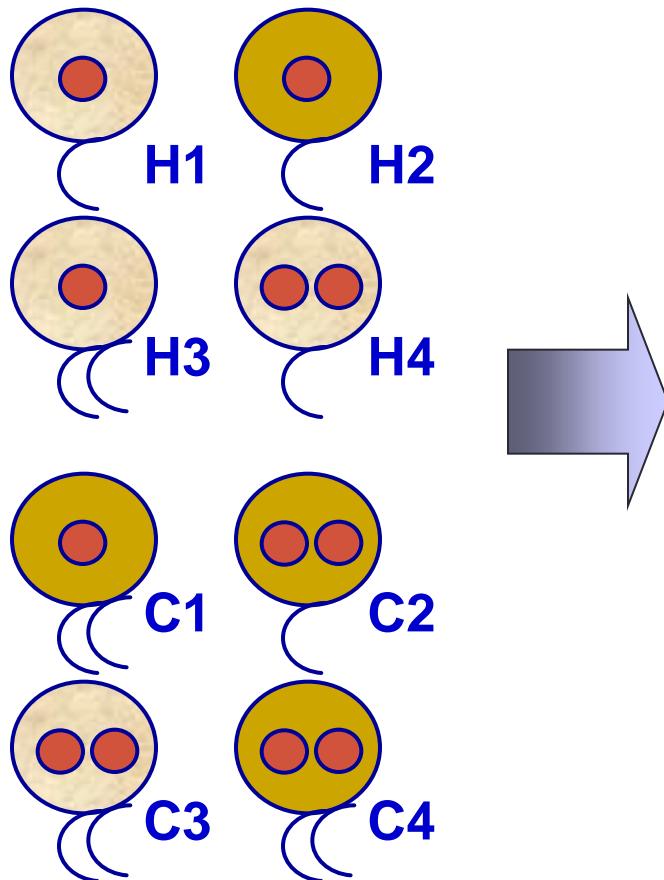
Dataset: Cancerous and Healthy Cells



	colour	#nuclei	#tails	class
H1	light	1	1	healthy
H2	dark	1	1	healthy
H3	light	1	2	healthy
H4	light	2	1	healthy
C1	dark	1	2	cancerous
C2	dark	2	1	cancerous
C3	light	2	2	cancerous
C4	dark	2	2	cancerous

# Decision Rules

## Mining with Decision Rules



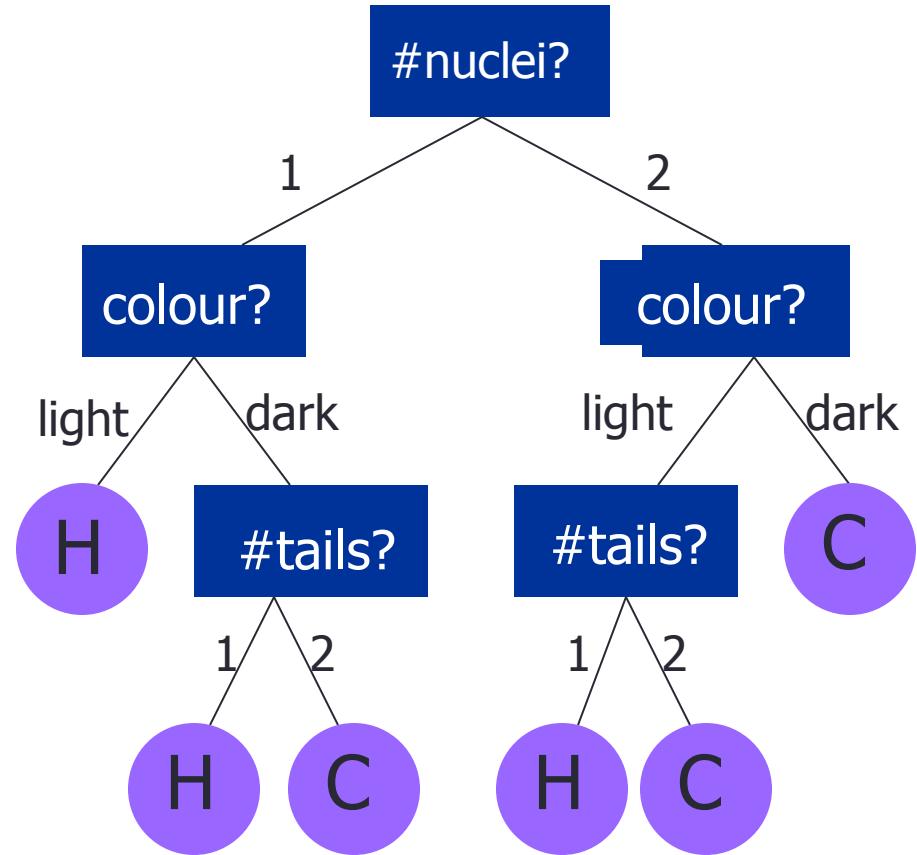
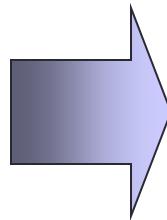
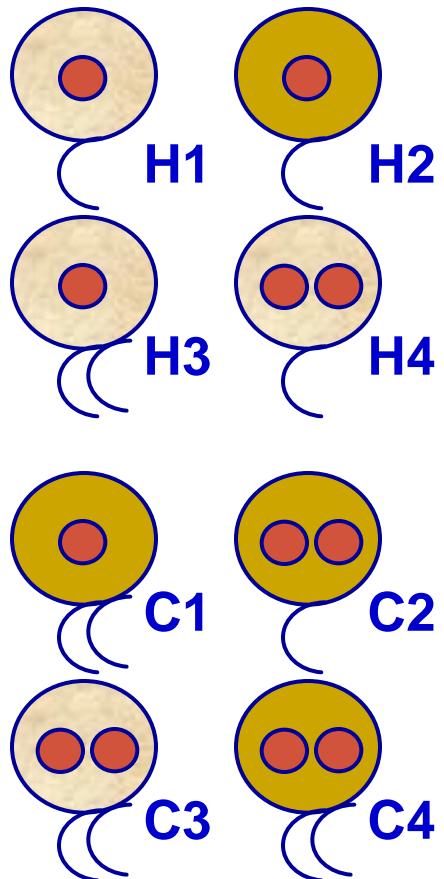
**If colour = light and # nuclei = 1  
Then cell = healthy**

**If #nuclei = 2 and colour = dark  
Then cell = cancerous**

**(and 4 rules more)**

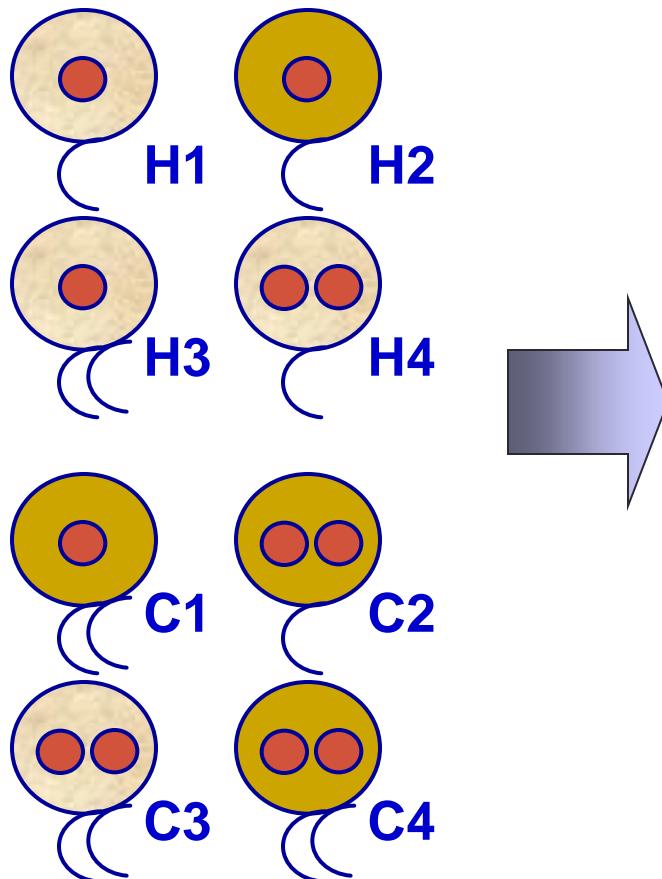
# Decision Trees

## Mining with Decision Trees



# Hierarchical Decision Rules

## Mining with Hierarchical Decision Rules



**If colour = light and # nuclei = 1  
Then cell = healthy**

**Else**

**If #nuclei = 2 and colour = dark  
Then cell = cancerous**

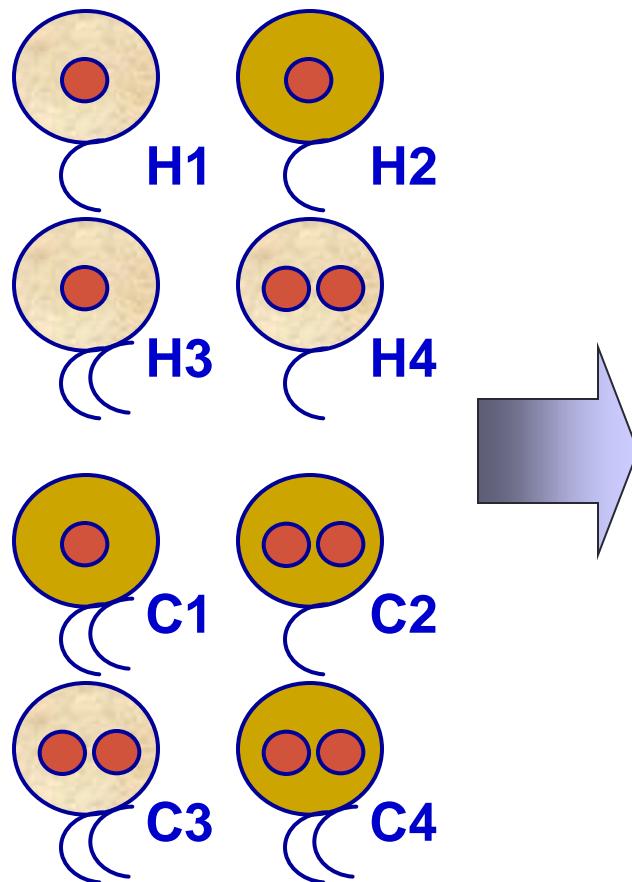
**Else**

**If #tails = 1  
Then cell = healthy**

**Else cell = cancerous**

# Association Rules

## Mining with Association Rules



If colour = light  
and # nuclei = 1

Then # tails = 1

(support = 25%;  
confidence = 50%)

If # nuclei = 2  
and cell = cancerous

Then # tails = 2

(support = 3/8%;  
confidence = 2/3%)

**association**

among A and B means that the presence of A in a record implies the presence of B in the same record

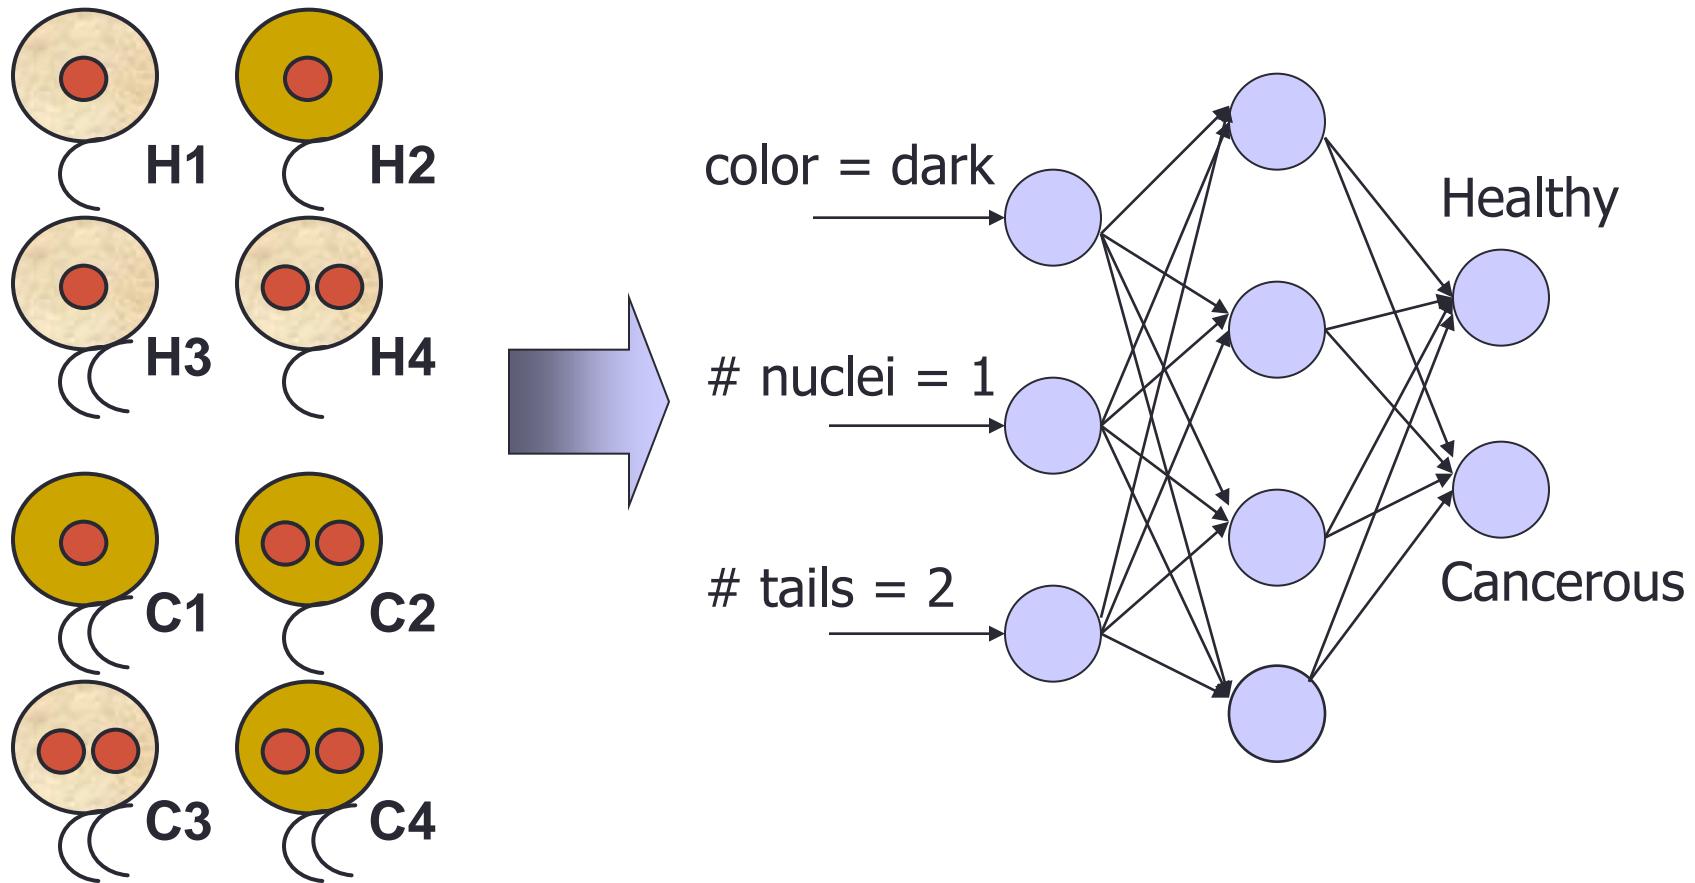
**Support:** the proportion of times that the rule applies.

**Confidence:** the proportion of times that the rule is correct

Apriori algorithm, Agrawal 1993

# Neural Networks

## Mining with Neural Networks

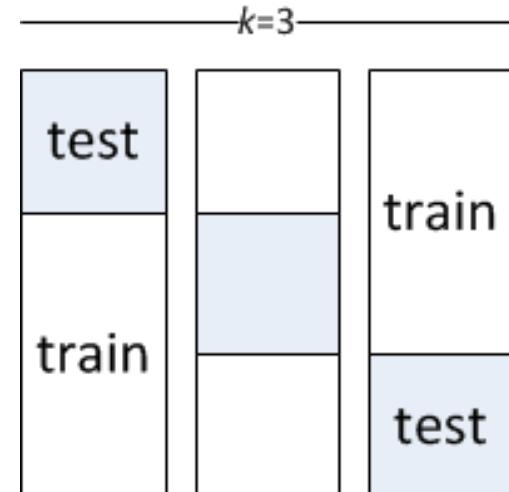


# Evaluation

- Once we obtain the model with the training data, we need to evaluate it with some new data (testing data)
  - We cannot use the same data for training and testing
    - E.g., Evaluating a student with the exercises previously solved
      - Student's marks will be "optimistic" and we don't know about student capability to generalise the learned concepts.
- Holdout approach consists of dividing the dataset into training (approx. 2/3 of the data) and testing (approx 1/3 of the data).
  - Problems: Data can be skewed, missing classes, etc. if randomly divided
- Stratification ensures that each class is represented with approximately equal proportions
  - E.g., if data contains approx 45% of positive cases, the training and testing datasets should maintain similar proportion of positive cases.
- Holdout estimate can be made more reliable by repeating the process with different subsamples (*repeated holdout* method)
  - The error rates on the different iterations are averaged (overall error rate)

# Cross-Validation

- *Cross-validation (CV)* avoids overlapping test sets
  - First step: split dataset ( $D$ ) into  $k$  subsets of equal size  $C_1, \dots, C_k$
  - Second step: we construct a dataset  $D_i = D - C_i$  used for training and test the accuracy of the classifier  $f_{D_i}$  on  $C_i$  subset for testing
    - Having done this for all  $k$  we estimate the accuracy of the method by averaging the accuracy over the  $k$  cross-validation trials
- Called *k-fold cross-validation*
  - Usually  $k=10$
  - Subsets are generally stratified before the CV is performed
  - The error estimates are averaged to yield an overall error estimate



# Evaluation Measures

- From the **confusion matrix**, we can obtain multiple measures about the goodness of a classifier.
- E.g. for a binary problem:

		Predicted		
		Positive	Negative	
Actual	Positive	<b>TP True Positive</b>	<b>FN False Negative (Type II error)</b>	$\text{TPRate} = \text{TP}/(\text{TP}+\text{FN})$ (Sensitivity, Recall)
	Negative	<b>FP False Positive (Type I error)</b>	<b>TN True Negative</b>	$\text{TNrate} = \text{TN}/(\text{FP}+\text{TN})$ (Specificity)
		$\text{PPV} = \text{TP}/(\text{TP}+\text{FP})$ Positive Predictive Value (Confidence, Precision)	$\text{NPV} = \text{TN}/(\text{FN}+\text{TN})$ Negative Predicted Value	Accuracy= $\text{TP}+\text{FP}/(\text{TP}+\text{TN}+\text{FP}+\text{FN})$

# Evaluation Measures

- Many times we need to combine the TP and FP to estimate the goodness of a classifier
  - For example, with imbalance data, the accuracy of a classifier needs to improve the percentage of the majority class. In a binay problem and 50/50 distribution, we need improve accuracy over 50%. However if the distribution is 90/10, accuracy needs to be over 90%
- *f – measure* is an harmonic median of these proportions:

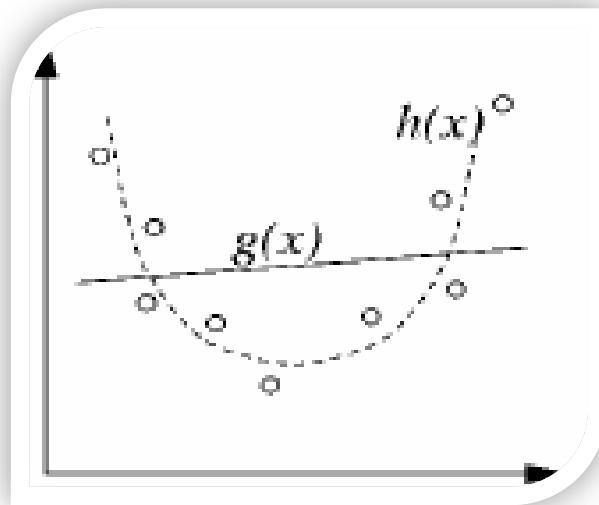
$$f - \text{measure} = \frac{2 \cdot TP}{2 \cdot TP + FP + FN}$$

- AUC (Area under the ROC)

# Evaluation: Underfitting vs. Overfitting

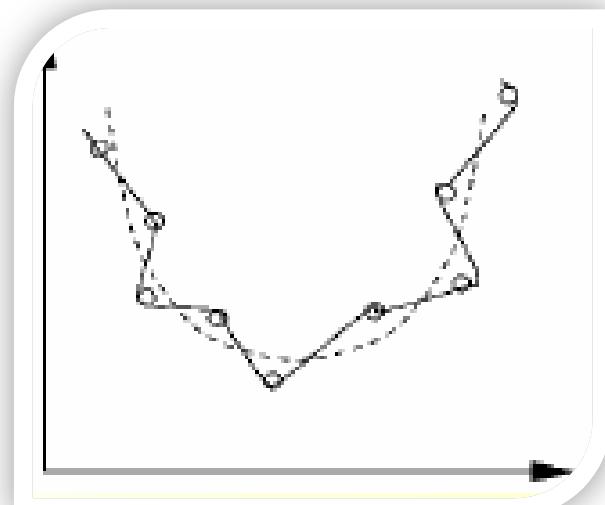
## Underfitting

- Too simple model



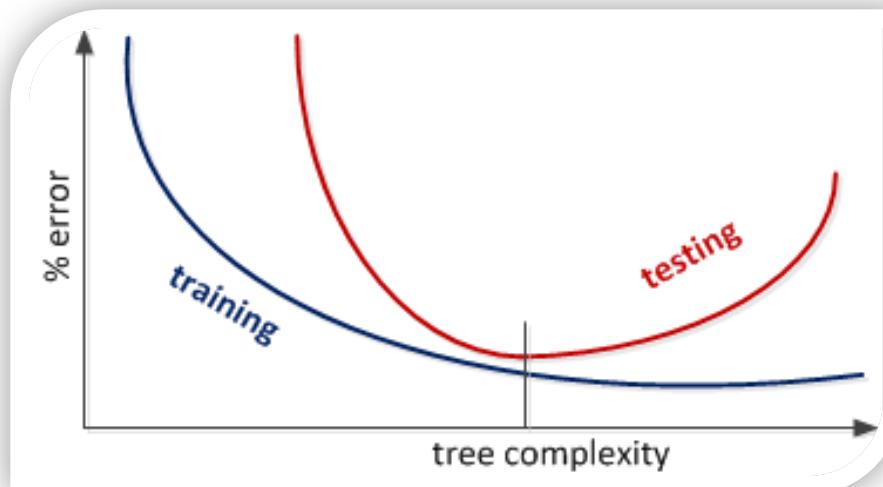
## Overfitting

- Too complex model



# Overfitting – Decision trees example

- Increasing the tree size, decreases the training and testing errors. However, at some point after (tree complexity), training error keeps decreasing but testing error increases
  - Many algorithms have parameters to determine the model complexity (e.g., in decision trees is the pruning parameter)



# Evaluation: Cost

- In many cases to fail in one way is not the same as failing in the other one (Type I error vs. Type II error).
- Cost can be considered when inducing the model from the data using the cost data.
  - Cost sensitive classifiers
  - Metacost (Dominos, 1999)
- Ej. Cost matrices
  - by default –left- and considering cost –right-:

		Predicted	
		Positive	Negative
Actual	Positive	0	1
	Negative	1	0

		Predicted	
		Positive	Negative
Actual	Positive	0	15
	Negative	1	0

# Thanks!

- Questions?