

# Managing Groups of Files in a Rule Oriented Data Management System (iRODS)

---

Andrea Weise

Reading University

---

June 23, 2008

Mike Wan

Wayne Schroeder

Adil Hasan



# Content

## 1 Introduction

## 2 SRB

## 3 iRODS

## 4 Mounted Collection

- Requirements
- Finding the Right System
- SQLite
- HAAW
- 7zip

## 5 Summery

## 6 References

# Definitions (1/2)

## Grid

Unites many systems within a virtual network

## Data Grid

Access to distributed data resources or repositories

## Data Grid Management System

Middleware software to “glue” resources together

# Definitions (1/2)

## Grid

Unites many systems within a virtual network

## Data Grid

Access to distributed data resources or repositories

## Data Grid Management System

Middleware software to “glue” resources together

# Definitions (1/2)

## Grid

Unites many systems within a virtual network

## Data Grid

Access to distributed data resources or repositories

## Data Grid Management System

Middleware software to “glue” resources together

# Definitions (2/2)

## SRB - Storage Resource Broker

- Data Grid middleware developed at the San Diego Supercomputer Center (SDSC)
- Distributed file system, based on a client-server architecture
- Access files seamlessly across a distributed environment

## iRODS - Rule Oriented Data System

- New Data Grid middleware developed at the San Diego Supercomputer Center (SDSC)
- Same features than SRB and more
- Express management policies as rules

# Definitions (2/2)

## SRB - Storage Resource Broker

- Data Grid middleware developed at the San Diego Supercomputer Center (SDSC)
- Distributed file system, based on a client-server architecture
- Access files seamlessly across a distributed environment

## iRODS - Rule Oriented Data System

- New Data Grid middleware developed at the San Diego Supercomputer Center (SDSC)
- Same features than SRB and more
- Express management policies as rules

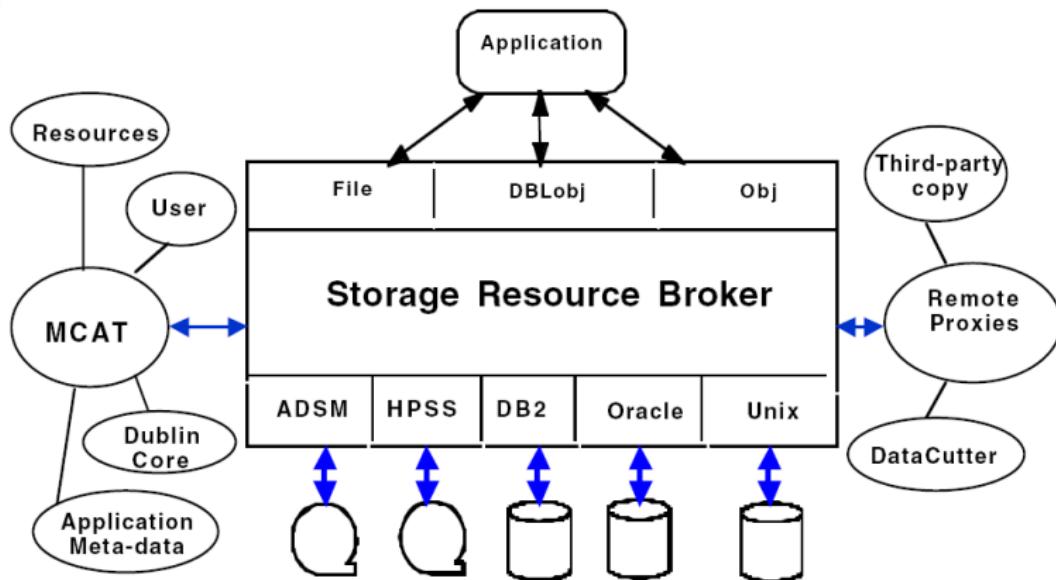
# SRB Overview

- Location transparency
- Improved reliability and availability
- Fault tolerance
- Single Sign On
- Integrated data access
- Federation

# SRB Acceptance

- Used in more than 100 projects world-wide
- More than a PetaByte stored in SRB systems
  - Taiwan National Data Archives
  - BIRN - Biomedical Informatics Research Network
  - National Optical Astronomers Observatory (NOAO)
  - UK eScience
  - ....

# Architecture



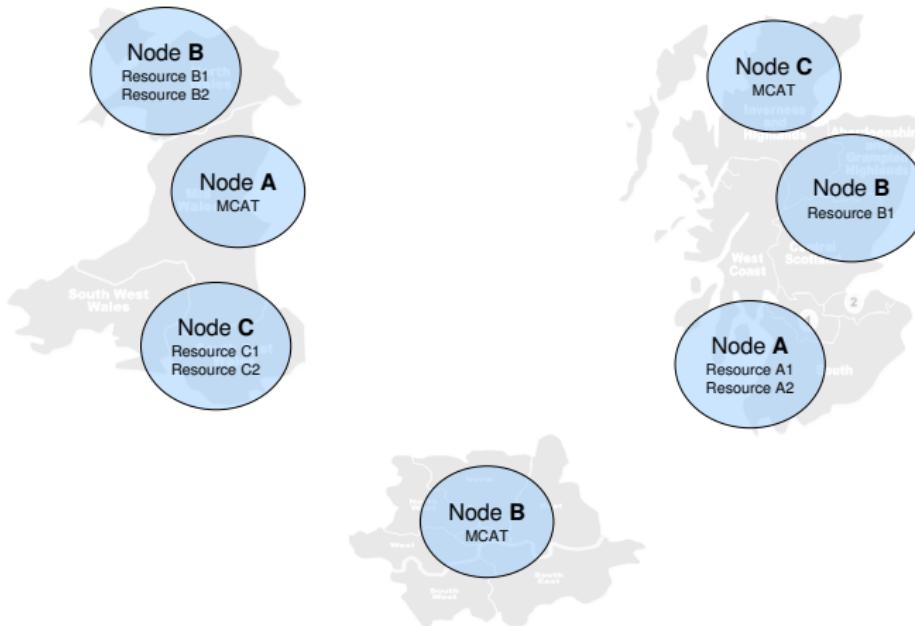
1

<sup>1</sup>R. Moore. Knowledge-based grids. IEEE Symposium on Mass Storage Systems and Technologies, IEEE Computer Society, 2003

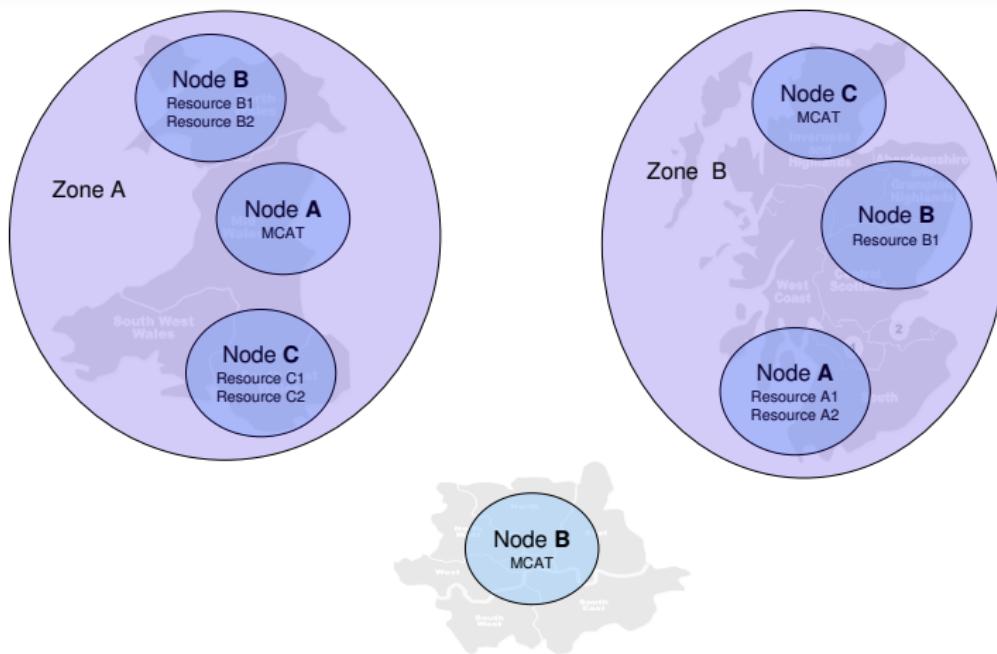
# Example



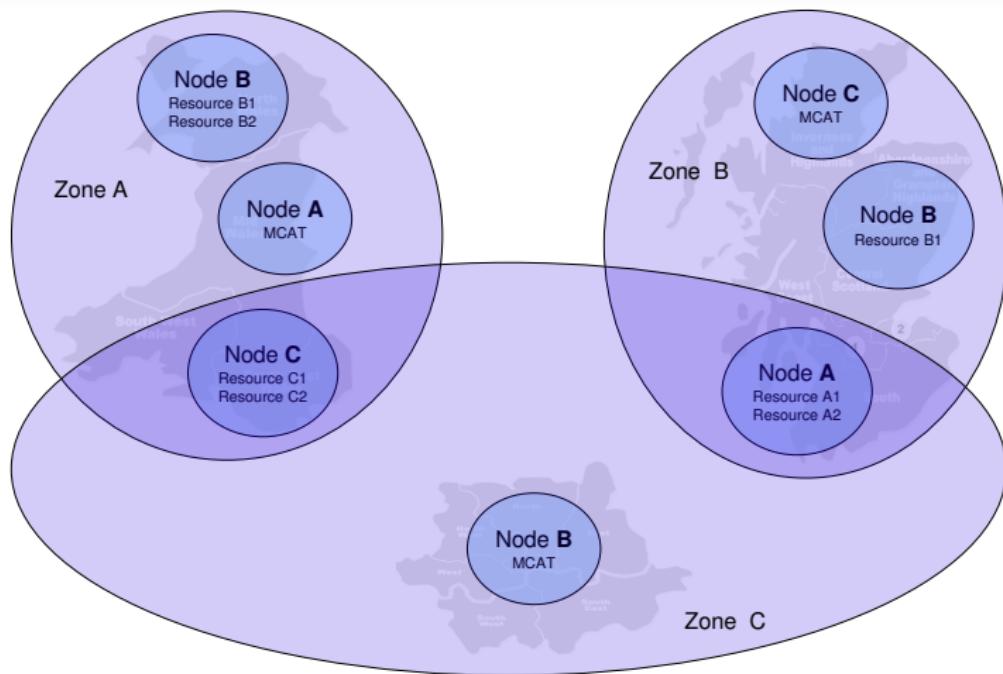
# Example



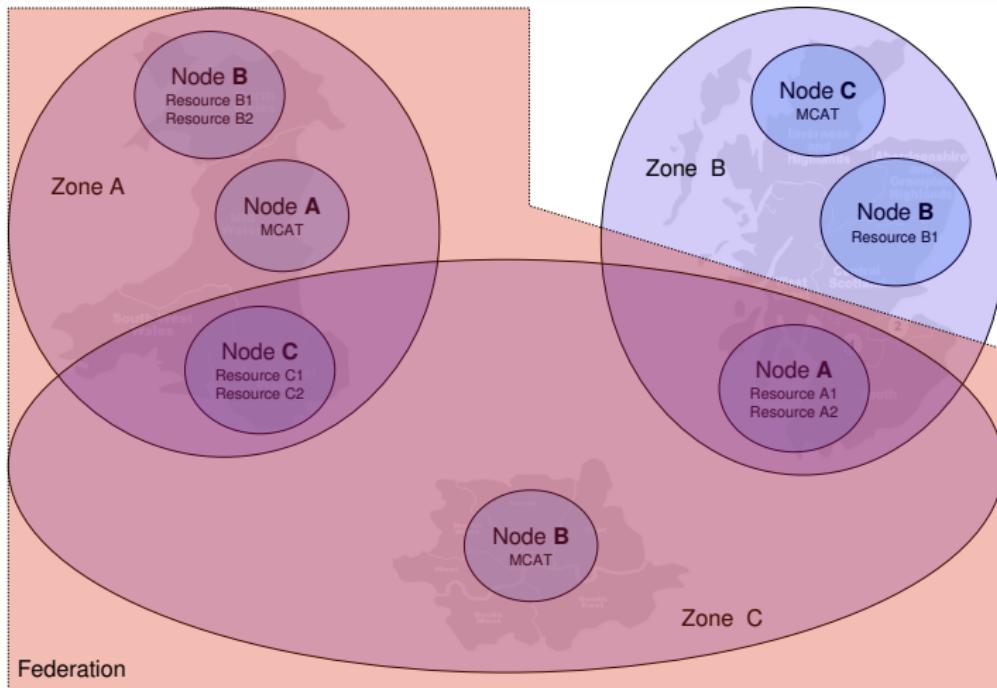
# Example



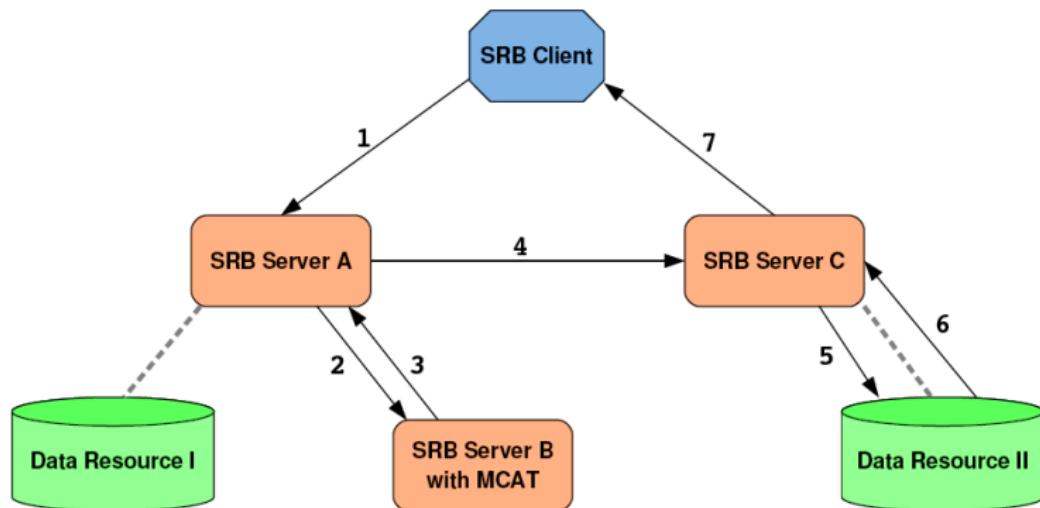
# Example



# Example



# Data Flow



2

<sup>2</sup>based on <http://www.vislab.uq.edu.au/research/SRB/background.html>

# Why a new system ?

## SRB Shortcomings

- SRB is quite complex
- Many operations
- Intelligence is hard-code
- Extensions/modifications require extreme care
- Does not allow incorporation of data management policies

## Aims

- System which is more flexible than SRB
  - Additional post processing on ingestion
  - Use workflows for server-side data management
  - Provide queued and batch processing
- Rule-based architecture to provide flexibility

# Why a new system ?

## SRB Shortcomings

- SRB is quite complex
- Many operations
- Intelligence is hard-code
- Extensions/modifications require extreme care
- Does not allow incorporation of data management policies

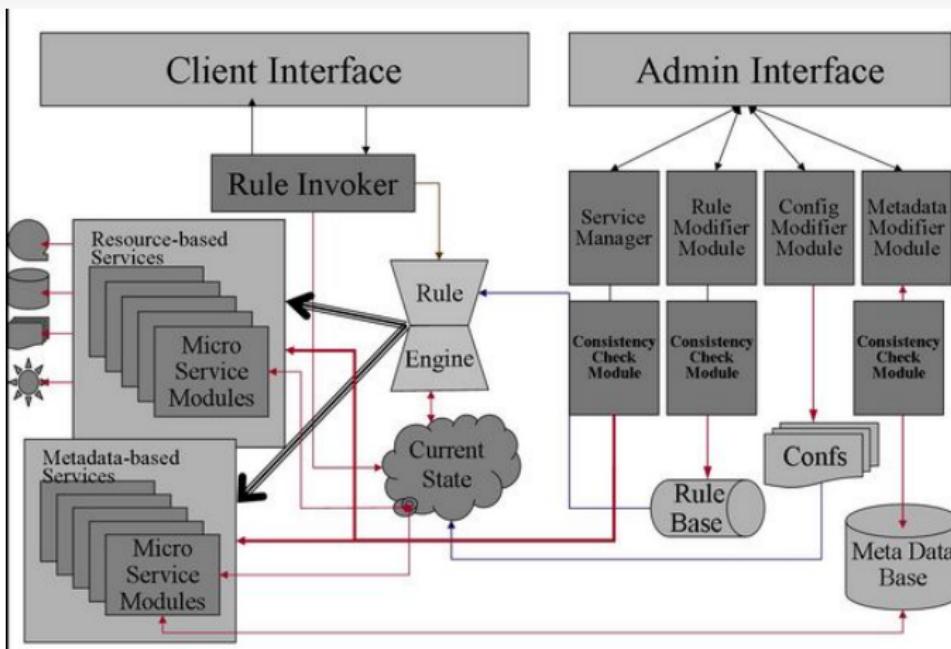
## Aims

- System which is more flexible than SRB
  - Additional post processing on ingestion
  - Use workflows for server-side data management
  - Provide queued and batch processing
- Rule-based architecture to provide flexibility

# iRODS Overview

- Open source client-server middleware
- Rule based management
- User can individualise work flow
- Only needed services compiled

# Architecture



# Rules

- Instructions, triggered by events
- Microservices (C-Functions)
- System individualisation
- Pre- and post processing

## Rule Format

actionDef | condition | workflow-chain | recovery-chain

## Example

HAAW(\*A,\*B) || HAAW-Bundle(\*A,\*B) | nop

# Rules

- Instructions, triggered by events
- Microservices (C-Functions)
- System individualisation
- Pre- and post processing

## Rule Format

actionDef | condition | workflow-chain | recovery-chain

## Example

HAAW(\*A,\*B) | | HAAW-Bundle(\*A,\*B) | nop

# Rules

- Instructions, triggered by events
- Microservices (C-Functions)
- System individualisation
- Pre- and post processing

## Rule Format

actionDef | condition | workflow-chain | recovery-chain

## Example

HAAW(\*A,\*B) || HAAW-Bundle(\*A,\*B) | nop

# Rule Types

- **Atomic** - applied immediately
- **Deferred** - support deferred consistent constraints
- **Periodic** - typically used to validate assertions

# Motivation

## *THE PROBLEM !*

- Small files - inefficient use of mass storage capabilities
- Improve usage of storage space
- Overcome performance problems within SRB - container concept
- Container <sub>downsides</sub>
  - No extensible metadata functionality
  - Not very efficient
  - No decrease in size
  - Tied to SRB
  - No container handling outside SRB

# Motivation

*WHAT DO WE WANT !*

- Replace SRB containers
- Archives within the system
- Archive handling in and outside iRODS
- Improve handling of huge amount of small files within iRODS
- Extensible metadata capabilities
- Plug- and Play

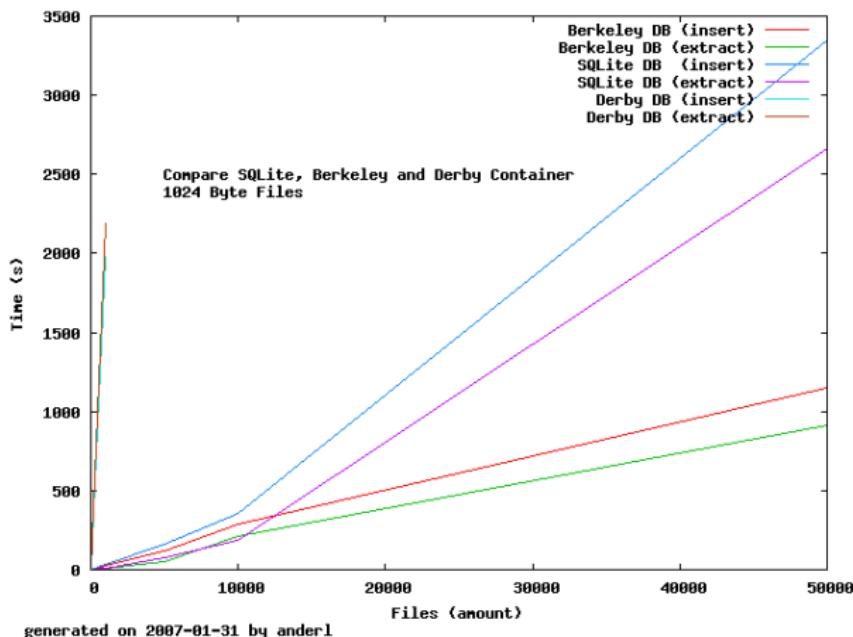
# Requirements

- Mounted Collections
- Basic requirements
  - Store all types of files
  - Add, delete, update files
  - Dynamically adjust of archive size
  - Accessable/usable in- and outside of iRODS
  - Searchable
  - Metadata
  - Single File only
  - Open Source
- Further requirements
  - Performance
  - Small Overhead compared to stored data
  - Compression
  - Encryption
  - Platform independent

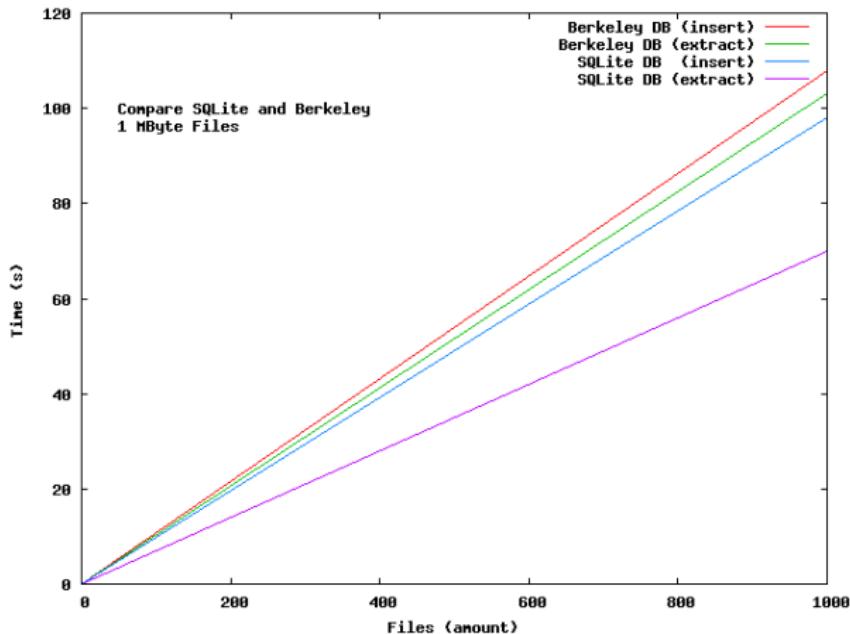
# Finding the Right System

- Database engine
  - Structured
  - Fast
  - Safe
- Single file only database
- Open source engines
  - Berkeley Database
  - Apache Derby
  - SQLite
- Standalone application

# Database Performance Tests



# Database Performance Tests



# Why SQLite

- Derby
  - Java
  - Very slow
- Berkely
  - Not single database file although described as single file only database
  - Slower with larger files

# SQLite

*"Small. Fast. Reliable."*<sup>4</sup>

- Single-file Cross-platform Database
- Zero-Configuration
- Cross-platform (Linux (Unix), MacOSX, OS/2, Win32 and WinCE)
- Self-Contained
- Public Domain (free to copy, modify, publish, use, compile, sell, or distribute SQLite code)
- Transactions are ACID
  - **Atomic**
  - **Consistent**
  - **Isolated**
  - **Durable**

---

<sup>4</sup> <http://www.sqlite.org>

# HAAW Archive

- SQLite database
  - Performance
  - Open Source
  - Single File
  - Embedded
- Metadata capabilities
- Independent of iRODS
- Hybrid of 7zip and SQLite
  - Overcome limitations in space
    - terabyte-sized databases
    - maximum length of a string or BLOB - 2 GB
  - Compression algorithm (LZMA compression<sup>5</sup>)
  - Encryption algorithm (AES-256 encryption<sup>6</sup>)

---

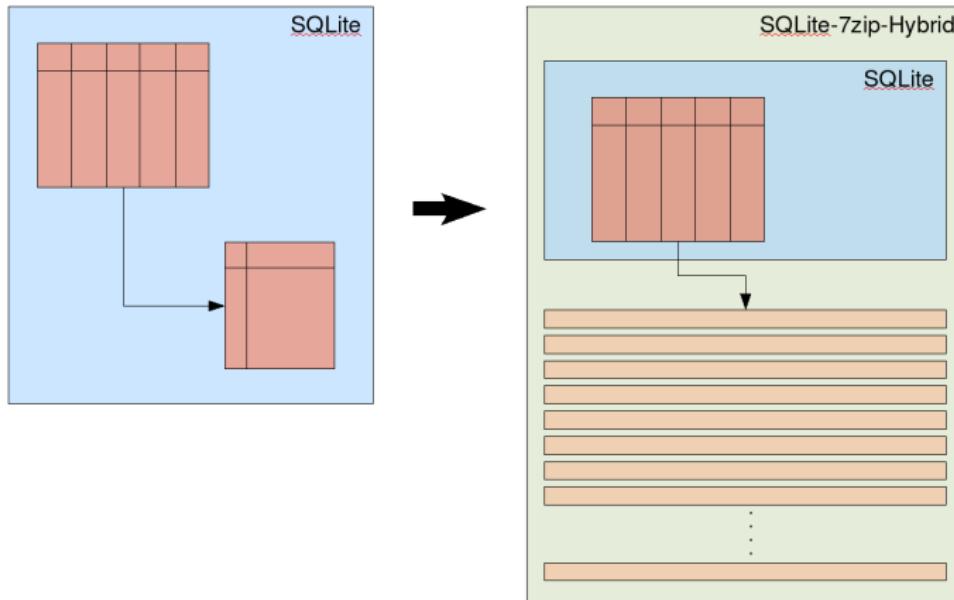
<sup>5</sup> LZMA - Lempel-Ziv-Markov chain-Algorithm

<sup>6</sup> AES - Advanced Encryption Standard

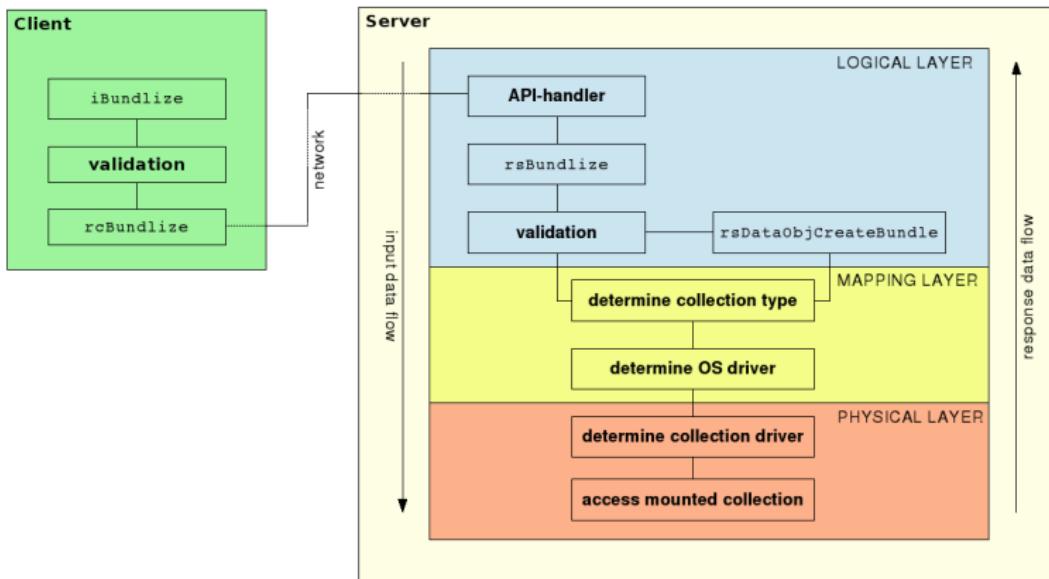
# 7zip

- File archiver
- High compression ratio
- Open source software (most source code under GNU LGPL license)
- LZMA SDK (Software Development Kit)

# SQLite - 7zip - Hybrid



# Framework



Framework mostly developed by Mike Wan (San Diego Supercomputer Center)

# Summary

- Definitions
- SRB
- iRODS
- Mounted Collections

# Thank you !

- Presentation Arcot (RAJA) Rajasekar, Data Grid Technologies Group, SDSC,UCSD, 2007
- R. Moore. Knowledge-based grids. IEEE Symposium on Mass Storage Systems and Technologies, IEEE Computer Society, 2003
- B. Strong, D. Corney, P. Berrisford, T. Folkes, C. Moreton-Smith, and K. Kleese- Van-Dam. Key lessons in the efficient archive of small files to the CCLRC MSS using SRB. Technical report, IEEE, 2005
- <http://irods.sdsc.edu/>
- <http://www.sqlite.org>
- <http://www.7-zip.org/>